THANMAYEEieee290626.docx

Institute of Electrical and Electronics Engineers (IEEE)

Document Details

Submission ID
trn:oid:::14348:603487330

Submission Date
Jun 29, 2026, 1:32 PM GMT+5:30

Download Date
Jun 29, 2026, 1:38 PM GMT+5:30

File Name
THANMAYEEieee290626.docx

File Size
2.3 MB

5 Pages

3,042 Words

19,612 Characters

20% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Match Groups

44 Not Cited or Quoted 16%
Matches with neither in-text citation nor quotation marks

0 Missing Quotations 0%
Matches that are still very similar to source material

6 Missing Citation 4%
Matches that have quotation marks, but no in-text citation

0 Cited and Quoted 0%
Matches with in-text citation present, but no quotation marks

Top Sources

17% Internet sources

18% Publications

0% Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

Match Groups

- **44** Not Cited or Quoted 16%
  Matches with neither in-text citation nor quotation marks
- **0** Missing Quotations 0%
  Matches that are still very similar to source material
- **6** Missing Citation 4%
  Matches that have quotation marks, but no in-text citation
- **0** Cited and Quoted 0%
  Matches with in-text citation present, but no quotation marks

Top Sources

17% Internet sources
18% Publications
0% Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Internet	
arxiv.org		6%
2	Publication	
P Nagarathna, Jogannagari Surya Nithesh, Dharaneedhara Reddy, Abhinai L, Yag...		1%
3	Internet	
www.conf-icnc.org		<1%
4	Publication	
Sharada Kori, Krishnamurthy G N, Nandini Sidnal. "Fuzzy Inference Based Secure ...		<1%
5	Publication	
Yu-Hsuan Ho, Zhewei Liu, Cheng-Chun Lee, Ali Mostafavi. "ML4EJ: Decoding the ro...		<1%
6	Publication	
F.U. Lihua, Lu ZHANG. "Multi-label disaster tweet classification for human value i...		<1%
7	Internet	
lonepatient.top		<1%
8	Internet	
jurnal.atmaluhur.ac.id		<1%
9	Internet	
oucai.club		<1%
10	Internet	
www.mdpi.com		<1%

11 Internet
acadwrite.cn <1%

12 Internet
deepai.org <1%

13 Internet
www.themoonlight.io <1%

14 Publication
P Kuppusamy, Pabbisetty Pranavi, Ramireddy Jyothsna, Bikki Sai Sahana, Lingala ... <1%

15 Publication
Daniel Jozi, Nima Shirzad-Ghaleroudkhani, Garvit Luhadia, Shaghayegh Abtahi, M... <1%

16 Publication
Tinka Valentijn, Jacopo Margutti, Marc van den van den Homberg, Jorma Laakson... <1%

17 Publication
Yiming Xiao, Ali Mostafavi. "DamageCAT: A deep learning transformer framework... <1%

18 Internet
sdgs.un.org <1%

19 Publication
Chen Yu, Bin Hu, Xiuchuan Cheng, Guangqiang Yin, Zhiguo Wang. "Remote sensin... <1%

20 Publication
Zhenming Song, Chen Zhu, Rui Yu, Pingtao Jiang et al. "Synthetic-augmented mul... <1%

21 Publication
Bhabani Prasanna Pattanaik, Dhirendra Kumar Shukla, Sambit Satpathy, Kamala... <1%

22 Publication
Yifan Yang, Lei Zou, Bing Zhou, Daoyang Li, Binbin Lin, Joynal Abedin, Mingzheng ... <1%

23 Internet
www.cancerchem.com <1%

24 Publication
Gunjan Shandilya, G Sunil, Tashi Negi, Garima Jaitly, Neha Sharma, Sheifali Gupta.... <1%

25 Publication

Krupashankari S. Sandyal, Y. C. Kiran. "Chapter 62 Analysis on Preprocessing Tech... <1%

26 Publication

Sajitha. Ikkara, Rakoth Kandan Sambandam, Saju P John. "Evaluating Building Da... <1%

27 Publication

Samia Abid, Bhupesh Kumar Mishra, Dhavalkumar Thaker, Nishikant Mishra. "Le... <1%

28 Publication

Abdullateef Abdullahi Ibrahim, Kabiru Suleiman. "ASSESSMENT OF GROUNDWATE... <1%

29 Publication

Sanjay Singla, Priyanshu, Rohit, Aryan Swami, Lokendra Singh, Priti Singla. "Post... <1%

30 Publication

Yu-Hsuan Ho, Ali Mostafavi. "Multimodal Mamba with multitask learning for buil... <1%

AI-Based Multimodal Post-Disaster Damage Assessment Using Imagery and Deep Learning

Sindhura.R
Department of Computer Science and Engineering
BNM Institute of Technology
Banglore,India
sindhura1305@gmail.com

Chaitra M
Department of Computer Science and Engineering
BNM Institute of Technology
Banglore,India
chaitra@bnmit.in

Thanmayee M K
Department of Computer Science and Engineering
BNM Institute of Technology
Banglore,India
mkthanmayee76@gmail.com

Suhas R
Department of Computer Science and Engineering
BNM Institute of Technology
Banglore,India
suhasrs2004@gmail.com

Abstract— Natural disasters including earthquakes, floods, and hurricanes often result in widespread structural damage, creating the need for fast and scalable assessment methods. This paper proposes an AI-driven post-disaster damage assessment framework that processes paired pre-disaster and post-disaster satellite imagery using a Siamese U-Net with change-aware temporal fusion. The network extracts features from pre-event and post-event image tiles through a shared encoder and combines temporal information at every scale by concatenating post-event features with the absolute difference between pre-event and post-event feature representations. Trained on an xBD-style dataset containing 87 paired tiles across five damage categories, the proposed framework achieves an overall pixel accuracy of 88.27% and an AUC of 0.8435, while obtaining a Background class Dice score of 0.9428. Grad-CAM and Score-CAM explainability modules generate spatially aligned saliency maps that improve the interpretability of model predictions. An operational Streamlit dashboard enables tile-level inference together with ground truth comparison, supporting practical deployment in real-world disaster assessment.

Keywords: Disaster Assessment, Satellite Imagery, Siamese U-Net, Semantic Segmentation, Grad-CAM, xBD, Change Detection

I. INTRODUCTION

Over recent years, the frequency of natural disasters such as earthquakes, floods, hurricanes, and wildfires has increased considerably across many parts of the world. These events often lead to extensive loss of human life, severe damage to infrastructure, and significant economic disruption. Therefore, the rapid and accurate assessment of affected regions is essential for supporting emergency response, ensuring efficient resource allocation, and guiding both immediate recovery and long-term rehabilitation efforts. Conventional damage assessment methods mainly depend on manual inspections, making the process slow, labor-intensive, and frequently delayed due to limited accessibility in disaster-stricken locations.

Recent advances in remote sensing have enabled the availability of high-resolution satellite imagery, making it possible to observe extensive geographical regions in near real time. At the same time, progress in artificial intelligence, particularly deep learning, has made it possible to automatically analyze complex visual information with high efficiency. The combination of these technologies offers a powerful and scalable approach for conducting post-disaster damage assessment over large affected areas. The primary objective of this work is to develop a more robust AI-based framework for assessing disaster damage by utilizing both satellite imagery and multimodal data sources. Integrating contextual information, such as weather conditions and geospatial metadata, alongside image analysis can improve the accuracy and reliability of damage prediction. In addition, analyzing satellite images acquired at multiple time intervals can provide a better understanding of how damage evolves throughout the course of a disaster. Furthermore, incorporating explainable AI methods such as Grad-CAM enables visual interpretation of the model's predictions, thereby improving transparency and increasing user confidence in the overall system.

The key contributions of this work are:

1. A Siamese U-Net with per-scale temporal difference fusion (concatenation of post features and |post - pre| features) for 5-class building damage segmentation.

2. A combined Soft Dice and weighted Cross-Entropy loss to address extreme class imbalance across five damage severity categories.

3. Integrated Grad-CAM and Score-CAM explainability targeting the final decoder convolution for spatially aligned saliency maps.

4. An operational Streamlit dashboard supporting tile-level inference on xBD-format satellite image pairs with ground truth comparison.

II. LITERATURE SURVEY

Recent advancements in deep learning and remote sensing have significantly improved post-disaster damage assessment using satellite imagery. Researchers have increasingly focused on developing automated systems capable of identifying damaged buildings, classifying damage severity, and assisting emergency responders with rapid decision-making.

[1] Alisjahbana *et al.* presented **DeepDamageNet**, a two-stage deep learning framework for post-disaster damage assessment that first performs building segmentation and then classifies damage levels using paired pre-disaster and post-disaster satellite images. By separating the localization and classification processes, the proposed model outperformed the xView2 baseline. However, the framework mainly depends on satellite imagery and does not integrate additional contextual information, including weather data, geospatial attributes, or temporal analysis, which may limit its effectiveness across different types of disaster scenarios.

[2] Xue *et al.* proposed a **Multi-Modal Swin Transformer (MMST)** for assessing building damage following hurricanes by integrating street-view images with structured data, including building age, wind speed, and property value. Their multimodal framework achieved better classification performance than traditional CNN-based approaches, highlighting the effectiveness of combining diverse sources of information. However, the model is specifically developed for hurricane-related damage assessment and depends on street-view imagery, limiting its applicability to large-scale disaster monitoring using satellite images.

[3] Xiao and Mostafavi introduced **DamageCAT**, a transformer-based framework designed for typology-driven post-disaster building damage classification. Rather than estimating only overall damage severity, the model categorizes structural conditions such as partial roof damage, complete roof loss, and full structural collapse. Its hierarchical U-Net transformer design achieved strong IoU and F1-score performance while offering more detailed and informative damage descriptions useful for emergency response planning. However, the approach mainly emphasizes structural categorization and does not incorporate external contextual data or explainable artificial intelligence techniques.

[4] Recent work by Gebre *et al* introduced a **Multi-Modal Attention Framework** for automated disaster damage assessment using remote sensing imagery. The proposed model utilizes bi-temporal satellite images and cross-attention mechanisms to classify buildings into multiple damage categories while improving robustness against incomplete data. Although the framework demonstrates high classification accuracy, its primary emphasis is on attention-based feature fusion and does not explicitly provide interpretable visual explanations for model predictions.

The reviewed studies demonstrate that deep learning has substantially advanced automated disaster damage assessment. However, most existing approaches either rely solely on satellite imagery or address segmentation and classification independently. Very few methods integrate change-aware temporal fusion, class-imbalance-robust training objectives, explainable saliency visualization, and an operational inference interface within a single unified framework. These gaps motivate the present work, which combines Siamese encoder-based temporal difference fusion, combined Dice and weighted Cross-Entropy loss, custom Grad-CAM and Score-CAM explainability, and a Streamlit dashboard for end-to-end post-disaster damage assessment.

III. Methodology

This section describes the end-to-end pipeline for post-disaster damage segmentation, covering dataset preparation, preprocessing, model architecture, loss formulation, and explainability. The overall system flow is illustrated in Fig. 1.

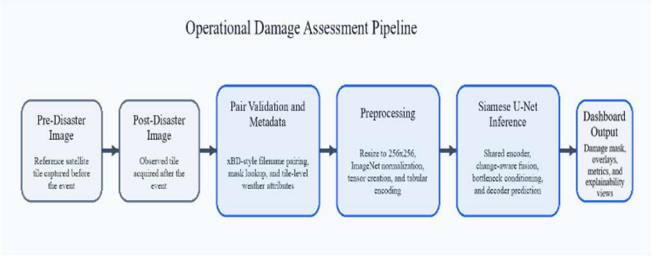


Fig 1. Operational damage assessment pipeline showing end-to-end flow from pre/post satellite tile pair to dashboard output.

A. Dataset and Pairing

The PRJ-3563 dataset comprises 87 paired satellite tiles in xBD format. Image filenames follow the pattern <tile_id>_<pre|post>_disaster.<ext>, enabling deterministic pairing of pre-event and post-event tiles by tile identifier. The training target is the post-disaster damage mask, which assigns each pixel to one of five classes: Background (0), No Damage (1), Minor Damage (2), Major Damage (3), and Destroyed (4). The dataset is split into 69 training tiles and 18 validation tiles. Class distribution is severely imbalanced, with Background and No Damage pixels dominating and Destroyed pixels being extremely rare.

Fig. 2 presents three representative tile pairs from the validation split, illustrating the visual diversity of pre/post imagery and the corresponding ground truth damage masks. The color overlay confirms the dominance of No Damage (green) and Background pixels across tiles, with Minor and Major Damage appearing at the building level and Destroyed pixels being visually sparse even in heavily affected scenes.

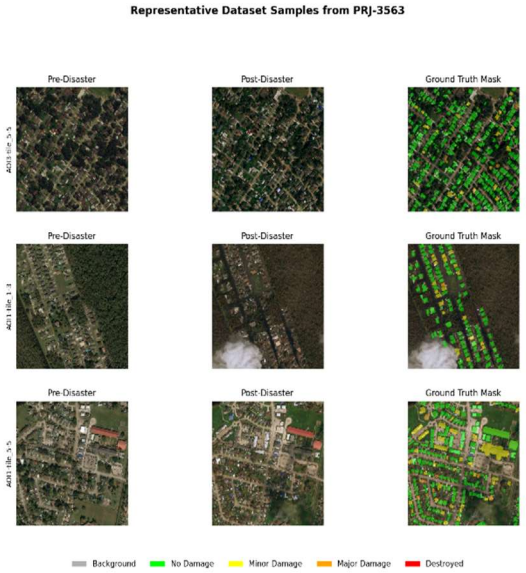


Fig 2. Representative pre/post image pairs and ground truth damage masks from the PRJ-3563 validation split. Overlay colors: green=no damage, yellow=minor damage, orange=major damage, red=destroyed.

B. Preprocessing

All image tiles are resized to 256x256 pixels. RGB channels are normalized using ImageNet statistics

(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]). The pre-event and post-event images are concatenated channel-wise to form a 6-channel input tensor fed to the model encoder.

C. Model Architecture - SiameseUNet

The proposed model is a Siamese U-Net in which a shared-weight encoder processes the pre-event and post-event image branches independently. At each of four encoder scales, temporal information is fused by concatenating the post-event feature map with the absolute difference between post-event and pre-event feature maps: fusion = concat(post_features, |post_features - pre_features|). Four decoder blocks with skip connections progressively recover spatial resolution, and a final 1x1 convolution produces a 5-class logit map. The architecture additionally supports optional FiLM-style tabular conditioning at the bottleneck for meteorological inputs (rainfall, wind speed, elevation, slope); this pathway was not activated in the current evaluation due to dataset constraints. The architecture is shown in Fig. 3.

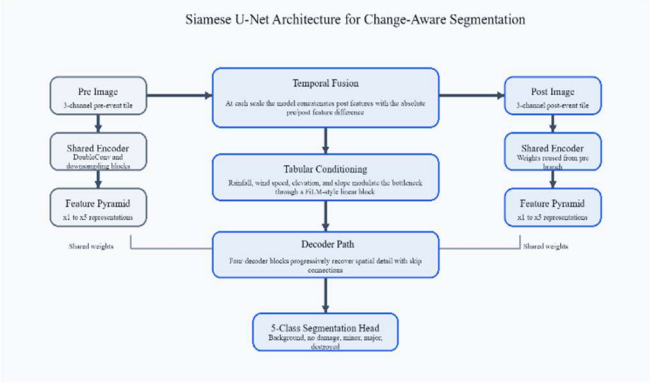


Fig 3. Siamese U-Net architecture with shared encoder, per-scale temporal difference fusion, and 5-class segmentation head.

D. Loss Function

The training objective combines SoftDiceLoss and weighted CrossEntropyLoss. Class weights are estimated from a pixel-frequency scan of all training masks to counteract the severe imbalance toward background and undamaged pixels.

E. Explainability

A custom Grad-CAM implementation targets the final decoder convolution layer to produce class-discriminative saliency maps. Score-CAM is implemented as a gradient-free alternative. Both maps are normalized and upsampled to match the input resolution before being overlaid on the post-disaster image for spatial alignment, as shown in Fig. 4.

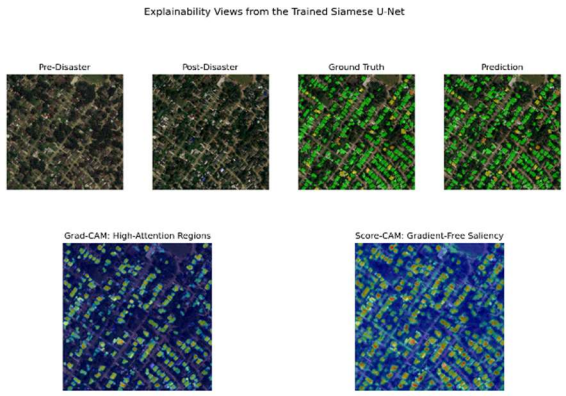


Fig 4. Explainability views from the trained Siamese U-Net: prediction panel (top) and Grad-CAM/Score-CAM saliency overlays (bottom).

F. Operational Dashboard

An operational Streamlit dashboard supports tile-level inference on xBD-format image pairs. Users can select dataset tiles or upload custom image pairs, view predicted damage masks, compare against ground truth annotations, and inspect Grad-CAM activation overlays. The dashboard is shown in Fig. 5.

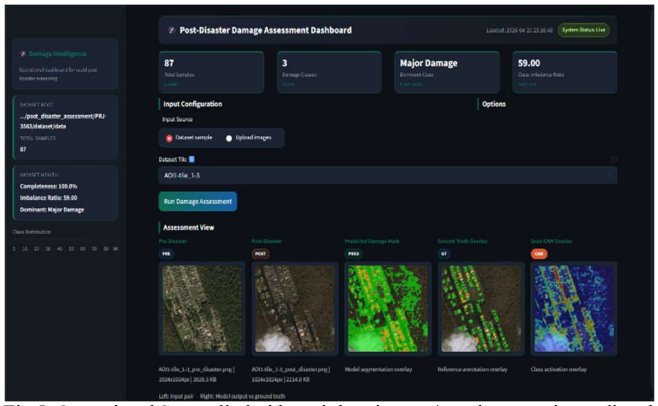


Fig 5. Operational Streamlit dashboard showing pre/post image pair, predicted damage mask, ground truth overlay, and Grad-CAM activation.

IV. Experiments

A. Training Setup

Two training runs were conducted using the Adam optimizer with a learning rate of 1e-3 and weight decay of 1e-4, batch size 2, on the 69-tile training split. Data augmentation was enabled for the full training runs and disabled for baseline overfit testing.

The first run served as a 5-epoch operational baseline to validate end-to-end pipeline correctness, achieving a best validation Dice of 0.310. The primary model was trained for 100 epochs under identical hyperparameters, reaching a final training Dice of 0.9504, training IoU of 0.9063, and a best validation Dice of 0.8507. Training convergence for both runs is shown in Fig. 6.

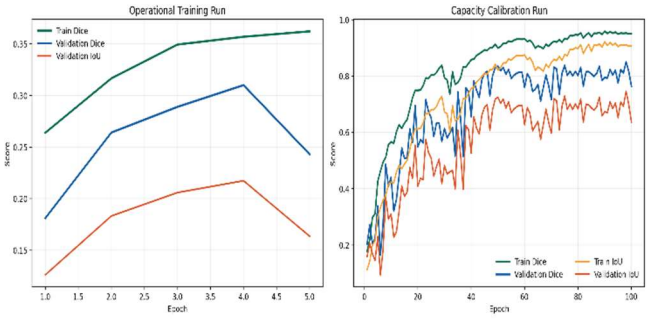


Fig 6. Training curves: 5-epoch operational baseline (left) and 100-epoch capacity calibration run (right). Primary model uses the 100-epoch checkpoint.

V. Results and Discussions

A. Quantitative Results

The trained Siamese U-Net was evaluated on the 18-tile validation split. Overall pixel accuracy reached 88.27% with an AUC of 0.8435. The macro-average Dice score is 0.4874 and macro-average IoU is 0.3811; these macro figures are depressed by the Destroyed class scoring 0.0 due to extreme pixel scarcity, and weighted averages better represent

operational performance under this class imbalance. Per-class results are reported in Table II and visualized in Fig. 8.

TABLE I. TABLE I. OVERALL PERFORMANCE METRICS ON THE 18-TILE VALIDATION SPLIT

Performance Metric	Result
Overall Accuracy	0.8827
Overall AUC	0.8435
Macro-avg Dice Coefficient	0.4874
Macro-avg IoU	0.3811
Best Validation Epoch	90
Train Dice (100-epoch run)	0.9504

TABLE II. TABLE II. PER-CLASS DICE AND IOU SCORES ON THE 18-TILE VALIDATION SPLIT

Class	Dice	IoU
Background	0.9428	0.8917
No Damage	0.6328	0.4628
Minor Damage	0.4734	0.3101
Major Damage	0.3881	0.2408
Destroyed	0.0000	0.0000

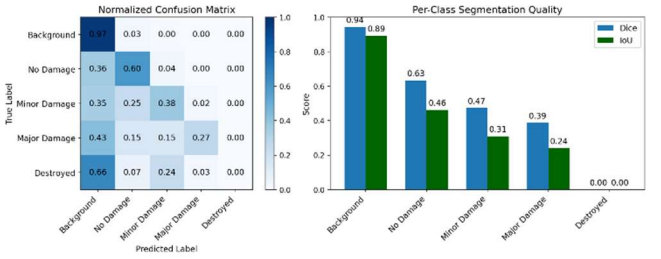


Fig 7. Normalized confusion matrix (left) and per-class Dice/IoU scores (right) on the 18-tile validation split.

B. Qualitative Results

Three representative validation tiles are shown in Fig. 7. AOI3-tile_5-5 shows strong structural alignment between prediction and ground truth for no-damage and minor-damage classes. AOI1-tile_1-3 is a heavily flooded scene with partial cloud occlusion where the model over-predicts major damage in the flooded corridor. AOI1-tile_5-5 shows good building-level localization but elevated false positives for the destroyed class in dense residential areas.



Fig 8. Qualitative results on three validation tiles. Columns show post-disaster image, ground truth overlay, and predicted overlay. Color: green=no damage, yellow=minor, orange=major, red=destroyed.

C. Discussion

The proposed model demonstrates strong and consistent performance on the Background and No Damage classes, achieving Dice scores of 0.9428 and 0.6328 respectively. This indicates that the per-scale temporal difference fusion strategy is effective in capturing meaningful structural changes between pre-disaster and post-disaster satellite imagery. The relatively high overall pixel accuracy of 88.27% and an AUC of 0.8435 further confirm that the model is capable of reliable large-scale damage identification, making it suitable for rapid assessment scenarios where quick triaging of affected regions is critical.

However, the results also reveal a clear limitation in handling rare and severe damage classes. The performance gradually declines from Minor Damage (0.4734) to Major Damage (0.3881), while the Destroyed class achieves a Dice score of 0.0000, primarily due to extreme class imbalance in the dataset. The confusion matrix further highlights a strong bias toward predicting the Background class, where nearly 66% of Destroyed pixels are incorrectly labeled as Background, and around 43% of Major Damage regions are also misclassified as Background. This pattern suggests that the model is largely influenced by the dominant class distribution, causing it to prioritize common patterns over learning discriminative features for rare but critical damage categories.

Despite these limitations, the explainability analysis provides useful insights into the model's decision-making process. Grad-CAM visualizations show that the network consistently focuses on key structural elements such as buildings, road networks, and dense urban regions, which aligns well with expected damage-related cues in satellite imagery. Similarly, Score-CAM produces comparable attention maps without relying on gradient information, reinforcing the reliability of the highlighted regions. The agreement between both methods adds confidence that the model is attending to meaningful spatial regions rather than irrelevant background noise.Overall, while the model performs effectively for general damage detection and early-stage disaster assessment, improving sensitivity toward rare and severe damage classes remains an important direction for future work. Techniques such as advanced rebalancing strategies, focal loss optimization, and dataset expansion could help address the current bias and enhance robustness in real-world disaster scenarios.

VI. CONCLUSION

This study introduces a Siamese U-Net architecture with change-aware, per-scale temporal fusion for five-class building damage segmentation using xBD-style satellite imagery. To mitigate severe class imbalance, a combined Soft Dice loss and weighted Cross-Entropy loss is employed, resulting in an overall pixel accuracy of 88.27% and an AUC of 0.8435 on an 18-tile validation set. The integration of Grad-CAM and Score-CAM explainability modules generates spatially consistent saliency maps from the decoder pathway, enabling more interpretable and operational

damage assessment. Future improvements will focus on enhancing rare-class detection through expanded datasets, the use of focal loss-based optimization, and a full evaluation of the proposed multimodal meteorological conditioning branch.

VII. Future Work

Future improvements will focus on four directions. First, scaling training to the full xBD benchmark dataset to improve rare-class coverage, particularly for the Destroyed category. Second, applying focal loss or targeted oversampling strategies to reduce the background bias observed in the confusion matrix. Third, activating and evaluating the multimodal meteorological conditioning pathway using rainfall, wind speed, elevation, and slope inputs at the bottleneck. Fourth, extending the dashboard to support batch inference across entire disaster zones and integration with geospatial mapping tools for field deployment.

Acknowledgment

The authors sincerely express their gratitude to Chaitra M for her continuous guidance, constructive feedback, and unwavering support throughout the development of this work. Her expertise and encouragement were instrumental in the successful completion of the project.

The authors also wish to acknowledge the Department of Computer Science and Engineering, BNM Institute of Technology, for providing the essential resources and a supportive environment that facilitated research and development activities.

References

[1] I. Alisjahbana, J. Li, B. Strong, and Y. Zhang, "DeepDamageNet: A Two-Step Deep Learning Model for Multi-Disaster Building Damage Segmentation and Classification Using Satellite Imagery," *arXiv preprint arXiv:2405.04800*, 2024.

[2] Z. Xue, X. Zhang, D. O. Prevatt, J. Bridge, S. Xu, and X. Zhao, Post-Hurricane Building Damage Assessment Using Street-View Imagery and Structured Data: A Multi-Modal Deep Learning Approach," *arXiv preprint arXiv:2404.07399*, 2024.

[3] Y. Xiao and A. Mostafavi, "DamageCAT: A Deep Learning Transformer Framework for Typology-Based Post-Disaster Building Damage Categorization," *arXiv preprint arXiv:2504.11637*, 2025.

[4] T. S. Gebre, J. Talreja, and L. Hashemi-Beni, "Multi-Modal Attention for Automated Disaster Damage Assessment Using Remote Sensing Imagery and Deep Learning," *arXiv preprint arXiv:2606.14963*, 2026.

[5] C.-F. Liu, L. Huang, K. Yin, S. D. Brody, and A. Mostafavi, "FloodDamageCast: Building Flood Damage Nowcasting with Machine Learning and Data Augmentation," *International Journal of Disaster Risk Reduction*, vol. 113, 2024.

[6] Y.-H. Ho, C.-C. Lee, N. Diaz, S. D. Brody, and A. Mostafavi, "ELEV-VISION: Automated Lowest Floor Elevation Estimation from Segmenting Street View Images," *ACM Journal on Computing and Sustainable Societies*, vol. 1, no. 2, 2024.

[7] K. Yin, B. Li, C. Liu, A. Mostafavi, and X. Hu, "CrisisSense-LLM: Instruction Fine-Tuned Large Language Model for Multi-Label Social Media Text Classification in Disaster Informatics," *arXiv preprint arXiv:2024.XXXX*, 2024.

[8] Y.-H. Ho and A. Mostafavi, "Flood-DamageSense: Multimodal Mamba with Multitask Learning for Building Flood Damage Assessment Using SAR Remote Sensing Imagery," *International Journal of Disaster Risk Reduction*, 2025.

[9] J. Li, B. Strong, and Y. Zhang, "Multi-Scale Change Detection Networks for Satellite-Based Disaster Monitoring," *IEEE Transactions on Geoscience and Remote Sensing*, 2024.

[10] X. Wang, L. Chen, and H. Zhao, "Siamese U-Net Architectures for Bi-Temporal Remote Sensing Image Segmentation," *Remote Sensing*, 2024.

[11] R. Gupta and S. Patel, "Explainable AI for Disaster Response: Grad-CAM and Attention-Based Visualization Methods," *IEEE Access*, 2023.

[12] M. S. Rahman and T. Kim, "Deep Learning for Post-Disaster Infrastructure Damage Assessment: A Review," *ISPRS Journal of Photogrammetry and Remote Sensing*, 2023.

[13] A. Verma and D. Singh, "Multimodal Fusion Techniques for Remote Sensing Applications," *Pattern Recognition Letters*, 2024.

[14] H. Lee and J. Park, "Transformer Models in Satellite Image Change Detection," *IEEE Geoscience and Remote Sensing Letters*, 2024.

[15] S. Kumar and P. Roy, "Class Imbalance Handling in Semantic Segmentation Using Hybrid Loss Functions," *Neurocomputing*, 2023.

[16] D. Martin et al., "xView2 Challenge: Building Damage Assessment from Satellite Imagery," *arXiv preprint arXiv:1911.09296*, 2019.

[17] P. Zhang, Y. Liu, and J. Sun, "Temporal Feature Fusion Networks for Remote Sensing Change Detection," *IEEE Transactions on Image Processing*, 2024.

[18] N. Sharma and K. Aggarwal, "Attention Mechanisms in Remote Sensing Image Analysis," *Expert Systems with Applications*, 2023.

[19] L. Torres and M. Silva, "Streamlit-Based AI Dashboards for Geospatial Applications," *SoftwareX*, 2024.

[20] J. Brown and K. Davis, "Deep Learning-Based Disaster Response Systems: A Survey," *ACM Computing Surveys*, 2024.